

CP4WatsonAIOps V3.2

Demo Environment Installation with Ansible



©2021 Niklaus Hirt / IBM

! THIS IS WORK IN PROGRESS

Please drop me a note on Slack or by mail nikh@ch.ibm.com if you find glitches or problems.

Changes

Date	Description	Files
17.09.2021	First Draft	
20.09.2021	Turbonomic, Humio and Tooling	
20.09.2021	Roles	
21.09.2021	Improved robustness and checks	
22.09.2021	Corrected some bugs	Thanks Henning Sternkicker
24.09.2021	Corrected some bugs in the debug script	Thanks Philippe Thomas
06.10.2021	Resiliency and Usability	
16.10.2021	Added EventManager (NOI) Standalone Option	
20.10.2021	Added AWX Option	Open Source Ansible Tower
21.10.2021	Added ManagelQ Option	Open Source Cloudforms
26.10.2021	10_debug_install.sh script updated	Still work in progress
27.10.2021	New template structure	
10.11.2021	First version for GA 3.2	
11.11.2021	Small tweaks	Demo and training scripts detect AIMgr and EvtMgr Namespace automatically
11.11.2021	Corrected missing wget	Thanks Nate Malone
19.11.2021	Humio install	Based on operator
23.11.2021	Added Service Mesh/Istio	Optional
24.11.2021	Change NOI to EventManager	including Namespace
20.11.2021	First stable release	v1.0
07.12.2021	Minor fixes	thanks Wei Huang

Installation

1. [Prerequisites](#)
2. [Architecture](#)
3. [AI Manager Base Install](#)
4. [Configure Applications and Topology](#)
5. [Training](#)
6. [Slack integration](#)
7. [Service Now integration](#)
8. [Some Polishing](#)
9. [Demo the Solution](#)
10. [Troubleshooting](#)
11. [Uninstall CP4WAIOPS](#)
12. [EventManager Install](#)
 - [Configure EventManager](#)
 - [Configure Runbooks](#)
13. [Installing Turbonomic](#)
14. [Installing ELK \(optional\)](#)
15. [Installing Humio \(optional\)](#)
16. [Installing ServiceMest/Istio \(optional\)](#)
17. [Installing and configuring AWX/AnsibleTower \(optional\)](#)

! You can find a PDF version of this guide here: [PDF](#).

Introduction

This repository contains the scripts for installing a Watson AIOps demo environment with an Ansible based installer.

They have been ported over from the shell scripts here <https://github.ibm.com/NIKH/aiops-3.1>.

As of 3.2 and going forward I will only update the Ansible scripts in this repository.

This is provided **as-is**:

- I'm sure there are errors
- I'm sure it's not complete
- It clearly can be improved

! This has been tested for the new CP4WAIOps 3.2 release on OpenShift 4.7 and 4.8.

I have tested on ROKS 4.7 and 4.8 and Fyre 4.6 and the scripts run to completion.

! Then EventManager/NOI-->AI Manager Gateway is not working yet on ROKS

So please if you have any feedback contact me

- on Slack: Niklaus Hirt or
- by Mail: nikh@ch.ibm.com

1. Prerequisites

1.1 OpenShift requirements

I installed the demo in a ROKS environment.

You'll need:

- ROKS 4.8 (4.7 should work also)
- 5x worker nodes Flavor `b3c.16x64` (so 16 CPU / 64 GB)

You might get away with less if you don't install some components (Humio, Turbonomic,...)

1.2 Tooling

You need the following tools installed in order to follow through this guide:

- ansible
- oc (4.7 or greater)
- jq
- kubectl (Not needed anymore - replaced by `oc`)
- kafkacat (only for training and debugging)
- elasticsearch (only for training and debugging)
- IBM cloudctl (only for LDAP)

1.2.1 On Mac - Automated (preferred)

You can either run:

```
sudo ./13_install_prerequisites_mac.sh
```

1.2.1.1 On Mac - Manual

Or install them manually:

```
/bin/bash -c "$(curl -fsSL
https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
brew install ansible
brew install kafkacat
brew install node
brew install wget
npm install elasticdump -g
brew install jq

curl -L https://github.com/IBM/cloud-pak-cli/releases/latest/download/cloudctl-darwin-
amd64.tar.gz -o cloudctl-darwin-amd64.tar.gz
tar xfvz cloudctl-darwin-amd64.tar.gz
sudo mv cloudctl-darwin-amd64 /usr/local/bin/cloudctl
rm cloudctl-darwin-amd64.tar.gz
```

Get oc and kubectl (optional) from [here](#)

or use :

```
wget https://github.com/openshift/okd/releases/download/4.7.0-0.okd-2021-07-03-
190901/openshift-client-mac-4.7.0-0.okd-2021-07-03-190901.tar.gz -O oc.tar.gz
tar xfvz oc.tar.gz
sudo mv oc /usr/local/bin
sudo mv kubectl /usr/local/bin. (this is optional)
rm oc.tar.gz
rm README.md
```

I highly recommend installing the **k9s** tool :

```
wget https://github.com/derailed/k9s/releases/download/v0.24.15/k9s_Darwin_x86_64.tar.gz
tar xfvz k9s_Darwin_x86_64.tar.gz
sudo mv k9s /usr/local/bin
rm LICENSE
rm README.md
```

1.2.2 On Ubuntu Linux - Automated (preferred)

For Ubuntu you can either run (for other distros you're on your own, sorry):

```
sudo ./14_install_prerequisites_ubuntu.sh
```

1.2.2.1 On Ubuntu Linux - Manual

Or install them manually:

sed comes preinstalled

```
sudo apt-get install -y ansible
sudo apt-get install -y kafkacat
sudo apt-get install -y npm
sudo apt-get install -y jq
sudo npm install elasticdump -g

curl -L https://github.com/IBM/cloud-pak-cli/releases/latest/download/cloudctl-linux-
amd64.tar.gz -o cloudctl-linux-amd64.tar.gz
tar xfvz cloudctl-linux-amd64.tar.gz
sudo mv cloudctl-linux-amd64 /usr/local/bin/cloudctl
rm cloudctl-linux-amd64.tar.gz
```

Get oc and oc from [here](#)

or use :

```
wget https://github.com/openshift/okd/releases/download/4.7.0-0.okd-2021-07-03-
190901/openshift-client-linux-4.7.0-0.okd-2021-07-03-190901.tar.gz -O oc.tar.gz
tar xfvz oc.tar.gz
sudo mv oc /usr/local/bin
sudo mv kubectl /usr/local/bin
rm oc.tar.gz
rm README.md
```

I highly recommend installing the **k9s** tool :

```
wget https://github.com/derailed/k9s/releases/download/v0.24.15/k9s_Linux_x86_64.tar.gz
tar xfvz k9s_Linux_x86_64.tar.gz
sudo mv k9s /usr/local/bin
rm LICENSE
rm README.md
```

1.3 Get the scripts and code from GitHub

1.3.1 Clone the GitHub Repository (preferred)

And obviously you'll need to download this repository to use the scripts.

```
git clone https://<YOUR GIT TOKEN>@github.ibm.com/NIKH/aiops-install-ansible.git
```

You can create your GIT token [here](#).

1.3.1.1 Refresh the code from GitHub

Make sure you have the latest version:

```
git checkout origin/master -f | git checkout master -f | git pull origin master
```

Or create an alias for reuse:

```
alias gitrefresh='git checkout origin/master -f | git checkout master -f | git pull origin master'
```

1.3.2 Download the GitHub Repository in a ZIP (not preferred)

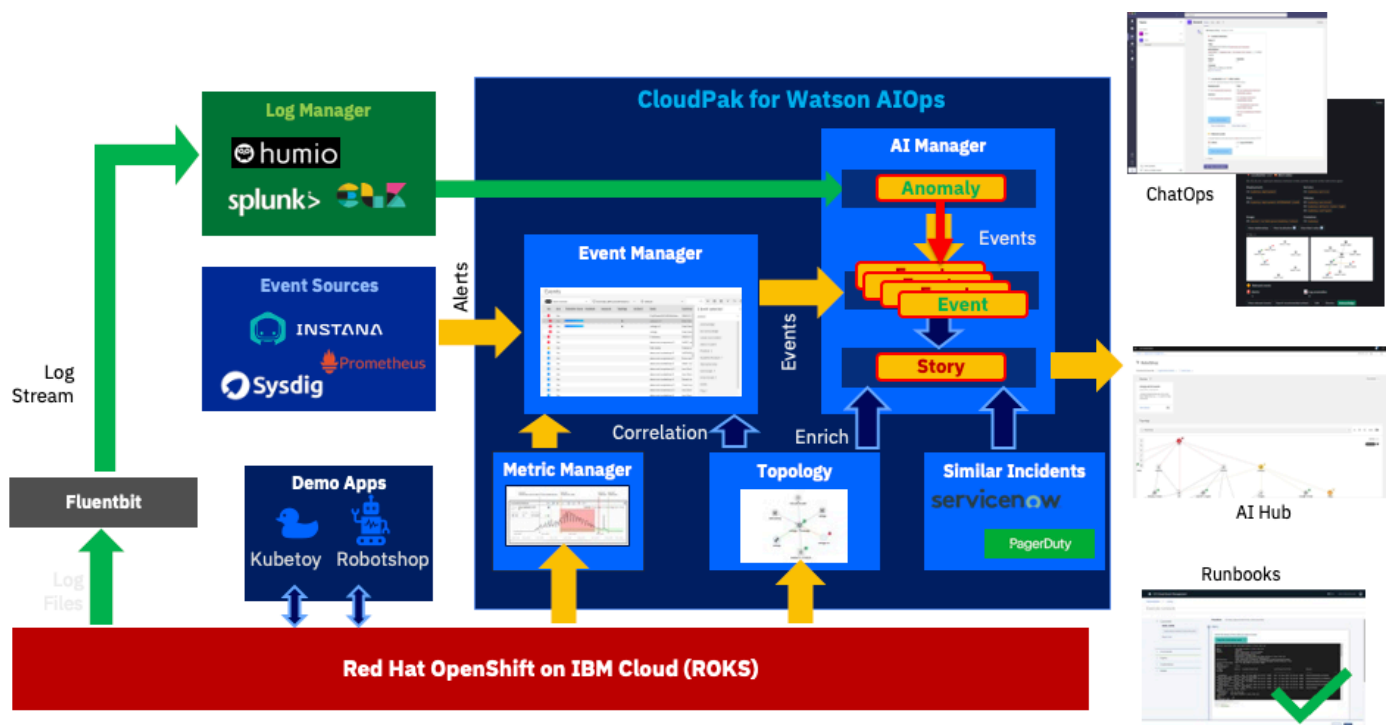
Simply click on the green **CODE** button and select **Download zip** to download the scripts and code.

! If there are updates you have to re-download the ZIP.

2. Architecture

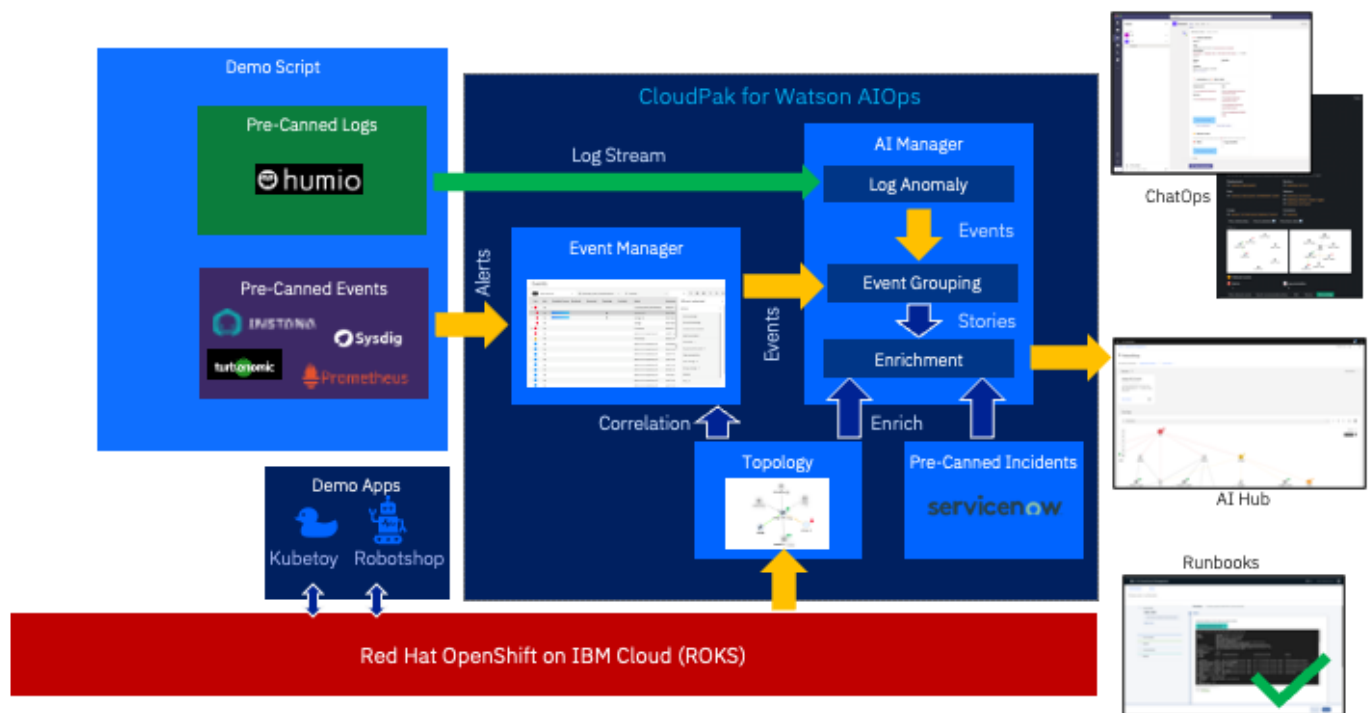
2.1 Basic Architecture

The environment (Kubernetes, Applications, ...) create logs that are being fed into a Log Management Tool (Humio in this case).



1. External Systems generate Alerts and send them into the EventManager (Netcool Operations Insight), which in turn sends them to the AI Manager for Event Grouping.
2. At the same time AI Manager ingests the raw logs coming from the Log Management Tool (Humio) and looks for anomalies in the stream based on the trained model.
3. If it finds an anomaly it forwards it to the Event Grouping as well.
4. Out of this, AI Manager creates a Story that is being enriched with Topology (Localization and Blast Radius) and with Similar Incidents that might help correct the problem.
5. The Story is then sent to Slack.
6. A Runbook is available to correct the problem but not launched automatically.

2.2 Optimized Demo Architecture



For the this specific Demo environment:

- Humio is not needed as I am using pre-canned logs for training and for the anomaly detection (inception)
- The Events are also created from pre-canned content that is injected into AI Manager
- There are also pre-canned ServiceNow Incidents if you don't want to do the live integration with SNOW
- The Webpages that are reachable from the Events are static and hosted on my GitHub
- The same goes for ServiceNow Incident pages if you don't integrate with live SNOW

This allows us to:

- Install the whole Demo Environment in a self-contained OCP Cluster
- Trigger the Anomalies reliably
- Get Events from sources that would normally not be available (Instana, Turbonomic, Metrics Manager, ...)
- Show some examples of SNOW integration without a live system

3. CP4WAIOPS Base Install

3.1 Install AI Manager

3.1.1 Adapt configuration

Adapt the `00_config_cp4waiops.yaml` file with the desired parameters:

3.1.1.1 Automatic Login

The Playbook provides the means to automatically login to the cluster by filling out the following section of the config file:

```
#
*****
*****
# -----
-----
# OCP LOGIN PARAMETERS
# -----
-----
#
*****
*****
OCP_LOGIN: true
OCP_URL: https://c100-e.eu-gb.containers.cloud.ibm.com:31513
OCP_TOKEN: sha256~T6-cxxxxxxxxxxxx-dtuj3ELQfpioUhHms

#Version of your OCP Cluster (override by setting manually - 4.6, 4.7,...)
OCP_MAJOR_VERSION: automatic
```

3.1.1.2 Adapt AI Manager Config

```
# #
*****
*****

# -----
-----

# CP4WAIOPS AI Manager INSTALL PARAMETERS
# -----
-----

#
*****
*****

# Install standalone AI Manager
INSTALL_AIOPS: true
# CP4WAIOPS Namespace for installation
WAIOPS_NAMESPACE: cp4waiops

# CP4WAIOPS Size of the install (small: PoC/Demo, tall: Production)
WAIOPS_SIZE: small # Leave at small unless you know what you're doing
# Version of the catalog subscription
SUBSCRIPTION_CHANNEL: v3.2

#
*****
*****

# -----
-----

# CP4WAIOPS Storage Class Override
# -----
-----

#
*****
*****

# Override the Storage Class auto detection (ibmc-file-gold-gid, rook-cephfs, nfs-
client, ...)
STORAGECLASS_FILE_OVERRIDE: not_configured
STORAGECLASS_BLOCK_OVERRIDE: not_configured
```

There is no need to manually define the Storage Class anymore.
The Playbook sets the storage class to `ibmc-file-gold-gid` for ROKS and `rook-cephfs` for Fyre.
Otherwise it uses the default Storage Class.

It is possible to override the Storage Class detection and force a custom Storage Class by setting `STORAGECLASS_XXX_OVERRIDE` in the config file.

3.1.1.3 Adapt Event Manager Config

```
#
*****
*****

# -----
-----
# CP4WAIOPS Event Manager INSTALL PARAMETERS
# -----
-----

#
*****
*****

# Install standalone Event Manager (Netcool Operations Insight)
INSTALL_EVTMGR: false
# CP4WAIOPS Namespace for installation
EVTMGR_NAMESPACE: cp4waiops-evtmgr
```

3.1.1.4 Adapt Demo Components

```
#
*****
*****

# -----
-----
# DEMO INSTALL PARAMETERS
# -----
-----

#
*****
*****

# Create a demo user in the OCP cluster
CREATE_DEMO_USER: true

# Install Demo Applications
INSTALL_DEMO_APPS: true

# Print all credentials at the end of the installation
PRINT_LOGINS: true

# Install Bastion Server for Runbook Automation
INSTALL_RUNBOOK_BASTION: true
```

3.1.1.4 Adapt Optional Components

```
#
*****
*****

# -----
-----

# MODULE INSTALL PARAMETERS
# -----
-----

#
*****
*****

# Install Rook-Ceph (Should Rook-Ceph be installed (automatic: install when on IBM Fyre)
(enable, automatic, disable))
ROOK_CEPH_INSTALL_MODE: automatic

# Install LDAP Server
INSTALL_LDAP: true

# Install Turbonomic (experimental - needs separate license)
INSTALL_TURBONOMIC: false
# Turbonomic Storage Class (ibmc-block-gold, rook-cephfs, nfs-client, ...)
STORAGE_CLASS_TURBO: ibmc-block-gold
# Install Turbonomic Metrics simulation (highly experimental!)
INSTALL_TURBONOMIC_METRICS: false
# Install Turbonomic --> Event Manager Gateway (highly experimental!)
INSTALL_TURBONOMIC_GATEWAY: false

# Install Humio (needs separate license)
INSTALL_HUMIO: false
# Humio Storage Class (ibmc-block-gold, rook-cephfs, nfs-client, ...)
STORAGE_CLASS_HUMIO: ibmc-block-gold

# Install ELK Stack
INSTALL_ELK: false

# Install ServiceMesh/Istio
INSTALL_ISTIO: false

# Install AWX (Open Source Ansible Tower)
INSTALL_AWX: false

# Install ManageIQ (Open Source CloudForms)
INSTALL_MANAGEIQ: false
```

3.1.2 Get the installation token

You can get the installation (pull) token from <https://myibm.ibm.com/products-services/containerlibrary>.

This allows the CP4WAIOPS images to be pulled from the IBM Container Registry.

This token is being referred to as <PULL_SECRET_TOKEN> below and should look something like this (this is NOT a valid token):

```
eyJhbGciOiJIUzI1NiJ9.eyJpc3adsGJJQk0gTWFYa2V0cGxhY2UiLCJpYXQiOiJlNzg0NzQzMjgsImp0aSI6IjRjYTM3gsdgdMzExNjQxZDdiMDJhMjRmMGMxMWgdsMzhIn0.Z-rqfSLJA-R-ow__tI3RmLx4mssdggdabvdcgdgYEkbYY
```


3.1.3 🚀 Start installation

Just run:

```
./10_install_ai_manager.sh -t <PULL_SECRET_TOKEN> [-v true]
```

Example:

```
./10_install_ai_manager.sh -t  
eyJhbGciOiJIUzI1NiJ9.eyJpc3adsGJJQk0gTWFYa2V0cGxhY2UiLCJpYXQiOiJlNzg0NzQzMjgsImp0aSI6IjR  
jYTM3gsdgdMzExNjQxZDdiMDJhMjRmMGMxMWgdsMzhIn0.Z-rqfSLJA-R-  
ow__tI3RmLx4mssdggdabvdcgdgYEkbYY
```

This will install:

- CP4WAIOPS AI Manager
- OpenLDAP (if enabled)
- Demo Apps (if enabled)
- Register LDAP Users (if enabled)
- Housekeeping
 - Additional Routes (Topology, Flink)
 - Create OCP User (serviceaccount demo-admin)
 - Patch Ingress
 - Adapt NGINX Certificates
 - Adapt Slack Welcome message to /welcome
- Turbonomic (if enabled)
- Humio (if enabled)
- OCP ELK Stack (if enabled)
- AWX (Open Source Ansible Tower - if enabled)
- ManagelQ (Open Source CloudForms - if enabled)

3.2 Install EventManager

To get the token, see [here](#)

3.2.1 🚀 Start installation

Just run:

```
./11_install_event_manager.sh -t <PULL_SECRET_TOKEN> [-v true]
```

Example:

```
./11_install_event_manager.sh -t  
eyJhbGciOiJIUzI1NiJ9.eyJpc3adsGJJQk0gTWFYa2V0cGxhY2UiLCJpYXQiOiJlNzg0NzQzMjgsImp0aSI6IjR  
jYTM3gsdgdMzExNjQxZDdiMDJhMjRmMGMxMWgdsMzhIn0.Z-rqfSLJA-R-  
ow__tI3RmLx4mssdggdabvdcgdgYEkbYY
```

This will install:

- CP4WAIOPS EventManager
- Gateway

3.3 Get Passwords and Credentials

At any moment you can run `./tools/20_get_logins.sh` that will print out all the relevant passwords and credentials.

Usually it's a good idea to store this in a file for later use:

```
./tools/20_get_logins.sh > my_credentials.txt
```

3.4 Check status of installation

At any moment you can run `./tools/11_check_install.sh` or for a more in-depth examination and troubleshooting `./tools/10_debug_install.sh` and select **Option 1** to check your installation.

4. Configure Applications and Topology

4.1 Create Kubernetes Observer for the Demo Applications

Do this for your applications (RobotShop by default)

- In the **AI Manager** "Hamburger" Menu select **Operate** / **Data and tool integrations**
- Click **Add connection**
- Under **Kubernetes**, click on **Add Integration**
- Click **Connect**
- Name it **RobotShop**
- Data Center **demo**
- Click **Next**
- Choose **local** for Connection Type
- Set **Hide pods that have been terminated** to **On**
- Set **Correlate analytics events on the namespace groups created by this job** to **On**
- Set Namespace to **robot-shop**
- Click **Next**
- Click **Done**

4.2 Create REST Observer to Load Topologies

- In the **AI Manager** "Hamburger" Menu select **Operate** / **Data and tool integrations**
- Click **Add connection**
- On the left click on **Topology**
- On the top right click on **You can also configure, schedule, and manage other observer jobs**
- Click on **Add a new Job**
- Select **REST** / **Configure**
- Choose "bulk_replace"
- Set Unique ID to "listenJob" (important!)
- Set Provider to whatever you like (usually I set it to "listenJob" as well)
- **Save**

4.3 Create Merge Rules for Kubernetes Observer

Launch the following:

```
./60_load_robotshop_topology_aimanager.sh
```

This will create:

- Merge Rules
- Merge Topologies for RobotShop.

! Please manually re-run the Kubernetes Observer to make sure that the merge has been done.

4.5 Create AIOps Application

Robotshop

- In the **AI Manager** go into **Operate** / **Application Management**
- Click **Create Application**
- Select **robot-shop** namespace
- Click **Add to Application**
- Name your Application (RobotShop)
- If you like check **Mark as favorite**
- Click **Save**

5. Training

5.1 Prepare Training

5.1.1 Create Kafka Humio Log Training Integration

- In the **AI Manager** "Hamburger" Menu select **Operate** / **Data and tool integrations**
- Click **Add connection**
- Under **Kafka**, click on **Add Integration**
- Click **Connect**
- Name it **HumioInject**
- Click **Next**
- Select **Data Source** / **Logs**
- Select **Mapping Type** / **Humio**
- Paste the following in **Mapping** (the default is **incorrect!**):

```
{
  "codec": "humio",
  "message_field": "@rawstring",
  "log_entity_types":
    "kubernetes.namespace_name,kubernetes.container_hash,kubernetes.host,kubernetes.container_name,kubernetes.pod_name",
  "instance_id_field": "kubernetes.container_name",
  "rolling_time": 10,
  "timestamp_field": "@timestamp"
}
```

- Click **Next**
- Toggle **Data Flow** to the **ON** position
- Select **Live data for continuous AI training and anomaly detection**
- Click **Save**

5.1.2 Create Kafka Netcool Training Integration

- In the **AI Manager** "Hamburger" Menu select **Operate** / **Data and tool integrations**
- Click **Add connection**
- Under **Kafka**, click on **Add Integration**
- Click **Connect**
- Name it **EvetnManager**
- Click **Next**
- Select **Data Source** / **Events**
- Select **Mapping Type** / **NOI**
- Click **Next**
- Toggle **Data Flow** to the **ON** position
- Click **Save**

5.1.3 Create ElasticSearch Port Forward

Please start port forward in **separate** terminal.

Run the following:

```
while true; do oc port-forward statefulset/iaf-system-elasticsearch-es-aiops 9200; done
```

or use the script that does it automatically

```
./tools/28_access_elastic.sh
```

5.2 Load Training Data

Run the following scripts to inject training data:

```
./50_load_robotshop_data.sh
```

This takes some time (20-60 minutes depending on your Internet speed).

5.3 Train Log Anomaly

5.3.1 Create Training Definition for Log Anomaly

- In the **AI Manager** "Hamburger" Menu select **Operate** / **AI model management**
- Under **Log anomaly detection - natural language** click on **Configure**
- Click **Next**
- Name it **LogAnomaly**
- Click **Next**
- Select **Custom**
- Select **05/05/21** (May 5th 2021 - dd/mm/yy) to **07/05/21** (May 7th 2021) as date range (this is when the logs we're going to inject have been created)
- Click **Next**
- Click **Next**
- Click **Create**

5.3.2 Train the Log Anomaly model

- Click on the **Manager** Tab
- Click on the **LogAnomaly** entry
- Click **Start Training**
- This will start a precheck that should tell you after a while that you are ready for training and then start the training

After successful training you should get:

The screenshot displays the Humio AI model management interface. The top navigation bar shows 'AI model management / Manage /'. The main header area includes 'Humio' and a 'Description' link. Below this, there are tabs for 'Overview', 'Versions', and 'Coverage'. The 'Overview' tab is active, showing a 'Training complete' status with a green checkmark and 'Models created'. The interface is divided into several sections: 'Schedule' (Training: Manual, Frequency: Not scheduled, At time: -), 'Data quality' (Good, 1 recommendation), 'Data set' (Name: Sm_7BHOB-I_dI7PlCfJU, Start date: 05/05/2021, End date: 05/07/2021), and 'Overview details' (AI type: Log anomaly detection - natural language, Version: v1, Version deployed: v1, Created on: 11/9/2021, 2:55:27 PM, Created by: admin). A 'Deploy' section at the bottom shows 'Deployment type: Manual' and 'Deployment date: 11/9/2021, 4:27:54 PM'.

- Click on **Deploy vXYZ**

⚠ If the training shows errors, please make sure that the date range of the training data is set to May 5th 2021 through May 7th 2021 (this is when the logs we're going to inject have been created)

5.4 Train Event Grouping

5.4.1 Create Training Definition for Event Grouping

- In the **AI Manager** "Hamburger" Menu select **Operate** / **AI model management**
- Under **Temporal grouping** click on **Configure**
- Click **Next**
- Name it **EventGrouping**
- Click **Next**
- Click **Done**

5.4.2 Train the Event Grouping Model

- Click on the **Manager** Tab
- Click on the **EventGrouping** entry
- Click **Start Training**
- This will start the training

After successful training you should get:

AI model management / Manage /

TemporalGrouping

Description ↗

Overview Versions Coverage

AI Training ⓘ

Training complete

Models created

Schedule ↗

Training Scheduled	Duration 11/9/2021 to 11/6/2025
Frequency Scheduled daily	Next scheduled job 11/11/2021
At time 12:00 AM	

Data quality ⓘ

Unavailable

Deploy

Deployment type When training is complete	Deployment date 11/10/2021, 12:14:53 AM
---	---

Start training ▶

Undeploy v2 ↗

Delete 🗑️

Overview details

AI type	Temporal grouping
Version	v2
Version deployed	v2
Created on	11/9/2021, 2:54:00 PM
Created by	admin

- The model is deployed automatically

5.5 Train Incident Similarity

! Only needed if you don't plan on doing the Service Now Integration

5.5.1 Create Training Definition

- In the **AI Manager** "Hamburger" Menu select **Operate** / **AI model management**
- Under **Similar incidents** click on **Configure**
- Click **Next**
- Name it **SimilarIncidents**
- Click **Next**
- Click **Next**
- Click **Done**

5.5.2 Train the Incident Similarity Model

- Click on the **Manager** Tab
- Click on the **SimilarIncidents** entry
- Click **Start Training**
- This will start the training

After successful training you should get:

AI model management / Manage /

SimilarIncidents
Description ↗

Overview Versions Coverage

AI Training ⓘ
Training complete
Models created

Schedule ↗

Training	Duration
Manual	-
Frequency	Next scheduled job
Not scheduled	Not scheduled
At time	-

Data quality ⓘ 1 recommendation
Good

Deploy

Deployment type	Deployment date
When training is complete	11/9/2021, 3:04:54 PM

Start training ▶
Undeploy v1 ↩
Delete 🗑

Overview details

AI type	Similar incidents
Version	v1
Version deployed	v1
Created on	11/9/2021, 2:54:24 PM
Created by	admin

- The model is deployed automatically

5.6 Train Change Risk

! Only needed if you don't plan on doing the Service Now Integration

5.6.1 Create Training Definition

- In the **AI Manager** "Hamburger" Menu select **Operate** / **AI model management**
- Under **Change risk** click on **Configure**
- Click **Next**
- Name it **ChangeRisk**
- Click **Next**
- Click **Next**
- Click **Done**

5.6.2 Train the Change Risk Model

- Click on the **Manager** Tab
- Click on the **ChangeRisk** entry
- Click **Start Training**
- This will start the training

After successful training you should get:

AI model management / Manage /

ChangeRisk
Description [🔗](#)

Overview Versions Coverage

AI Training ⓘ
Training complete
Models created

Schedule [🔗](#)

Training	Duration
Manual	-
Frequency	Next scheduled job
Not scheduled	Not scheduled
At time	-

Data quality ⓘ [1 recommendation](#)

Good

Start training [▶](#)

Undeploy v1 [🔗](#)

Delete [🗑](#)

Deploy [🔗](#)

Deployment type	Deployment date
Manual	11/9/2021, 3:06:25 PM

Overview details

AI type	Change risk
Version	v1
Version deployed	v1
Created on	11/9/2021, 2:55:52 PM
Created by	admin

- Click on **Deploy vXYZ**

6. Slack integration

6.1 Initial Slack Setup

For the system to work you need to setup your own secure gateway and slack workspace. It is suggested that you do this within the public slack so that you can invite the customer to the experience as well. It also makes it easier for is to release this image to Business partners

You will need to create your own workspace to connect to your instance of CP4WAOps.

Here are the steps to follow:

1. [Create Slack Workspace](#)
2. [Create Slack App](#)
3. [Create Slack Channels](#)
4. [Create Slack Integration](#)
5. [Get the Integration URL - Public Cloud - ROKS](#) OR
6. [Get the Integration URL - Private Cloud - Fyre/TEC](#)
7. [Create Slack App Communications](#)
8. [Prepare Slack Reset](#)

6.2 NGNIX Certificate for V3.1.1 - If the integration is not working

In order for Slack integration to work, there must be a signed certicate on the NGNIX pods. The default certificate is self-signed and Slack will not accept that. The method for updating the certificate has changed between AIOps v2.1 and V3.1.1. The NGNIX pods in V3.1.1 mount the certificate through a secret called `external-tls-secret` and that takes precedent over the certificates staged under `/user-home/_global_/customer-certs/`.

For customer deployments, it is required for the customer to provide their own signed certificates. An easy workaround for this is to use the Openshift certificate when deploying on ROKS. **Caveat:** The CA signed certificate used by Openshift is automatically cycled by ROKS (I think every 90 days), so you will need to repeat the below once the existing certificate is expired and possibly reconfigure Slack.

This method replaces the existing secret/certificate with the one that OpenShift ingress uses, not altering the NGINX deployment. An important note, these instructions are for configuring the certificate post-install. Best practice is to follow the installation instructions for configuring certificates during that time.

The custom resource `AutomationUIConfig/iaf-system` controls the certificates and the NGINX pods that use those certificates. Any direct update to the certificates or pods will eventually get overwritten, unless you first reconfigure `iaf-system`. It's a bit tricky post-install as you will have to recreate the `iaf-system` resource quickly after deleting it, or else the installation operator will recreate it. For this reason it's important to run all the commands one after the other. **Ensure that you are in the project for AIOps**, then

paste all the code on your command line to replace the `iaf-system` resource.

```
NAMESPACE=$(oc project -q)
IAF_STORAGE=$(oc get AutomationUIConfig -n $NAMESPACE -o jsonpath='{
.items[*].spec.storage.class }')
oc get -n $NAMESPACE AutomationUIConfig iaf-system -oyaml > iaf-system-backup.yaml
oc delete -n $NAMESPACE AutomationUIConfig iaf-system
cat <<EOF | oc apply -f -
apiVersion: core.automation.ibm.com/v1beta1
kind: AutomationUIConfig
metadata:
  name: iaf-system
  namespace: $NAMESPACE
spec:
  description: AutomationUIConfig for cp4waiops
  license:
    accept: true
  version: v1.0
  storage:
    class: $IAF_STORAGE
  tls:
    caSecret:
      key: ca.crt
      secretName: external-tls-secret
    certificateSecret:
      secretName: external-tls-secret
EOF
```

Again, **ensure that you are in the project for AIOps** and run the following to replace the existing secret with a secret containing the OpenShift ingress certificate.

```
WAIOPS_NAMESPACE=$(oc project -q)
# collect certificate from OpenShift ingress
ingress_pod=$(oc get secrets -n openshift-ingress | grep tls | grep -v router-metrics-
certs-default | awk '{print $1}')
oc get secret -n openshift-ingress -o 'go-template={{index .data "tls.crt"}}'
${ingress_pod} | base64 -d > cert.crt
oc get secret -n openshift-ingress -o 'go-template={{index .data "tls.key"}}'
${ingress_pod} | base64 -d > cert.key
oc get secret -n $WAIOPS_NAMESPACE iaf-system-automationui-aui-zen-ca -o 'go-template=
{{index .data "ca.crt"}}' | base64 -d > ca.crt
# backup existing secret
oc get secret -n $WAIOPS_NAMESPACE external-tls-secret -o yaml > external-tls-
secret$(date +%Y-%m-%dT%H:%M:%S).yaml
# delete existing secret
oc delete secret -n $WAIOPS_NAMESPACE external-tls-secret
# create new secret
oc create secret generic -n $WAIOPS_NAMESPACE external-tls-secret --from-
file=ca.crt=ca.crt --from-file=cert.crt=cert.crt --from-file=cert.key=cert.key --dry-
run=client -o yaml | oc apply -f -
#oc create secret generic -n $WAIOPS_NAMESPACE external-tls-secret --from-
file=cert.crt=cert.crt --from-file=cert.key=cert.key --dry-run=client -o yaml | oc apply
-f -
# scale down nginx
REPLICAS=2
oc scale Deployment/ibm-nginx --replicas=0
# scale up nginx
sleep 3
oc scale Deployment/ibm-nginx --replicas=${REPLICAS}
rm external-tls-secret
```

Wait for the nginx pods to come back up

```
oc get pods -l component=ibm-nginx
```

When the integration is running, remove the backup file

```
rm ./iaf-system-backup.yaml
```

And then restart the Slack integration Pod

```
oc delete pod $(oc get po -n $WAIOPS_NAMESPACE | grep slack | awk '{print $1}') -n
$WAIOPS_NAMESPACE --grace-period 0 --force
```

The last few lines scales down the NGINX pods and scales them back up. It takes about 3 minutes for the pods to fully come back up.

Once those pods have come back up, you can verify the certificate is secure by logging in to AIOps. Note that the login page is not part of AIOps, but rather part of Foundational Services. So you will have to login first and then check that the certificate is valid once logged in. If you want to update the certificate for Foundational Services you can find instructions [here](#).

6.3 Change the Slack Slash Welcome Message (optional)

If you want to change the welcome message

```
oc set env deployment/${(oc get deploy -l app.kubernetes.io/component=chatops-slack-  
integrator -o jsonpath='{.items[*].metadata.name }')} SLACK_WELCOME_COMMAND_NAME=/aiops-  
help
```

7. Service Now integration

7.1 Integration

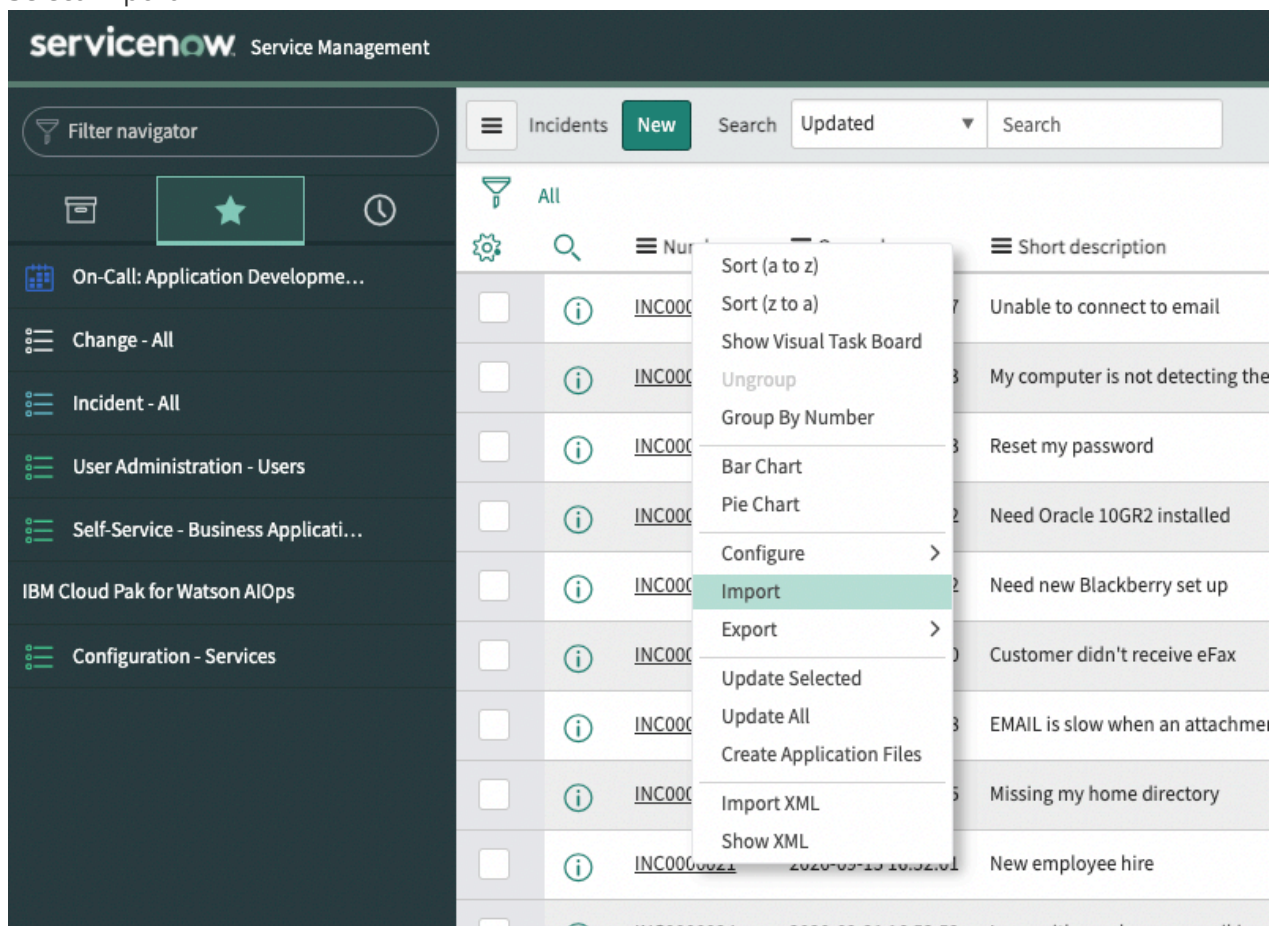
1. Follow [this](#) document to get and configure your Service Now Dev instance with CP4WAIOPS.

Stop at **Testing the ServiceNow Integration**.

!! Don't do the training as of yet.

2. Import the Changes from ./doc/servicenow/import_change.xlsx

1. Select **Change - All** from the right-hand menu
2. Right Click on **Number** in the header column
3. Select Import



4. Chose the ./doc/servicenow/import_change.xlsx file and click **Upload**

servicenow Service Management

Filter navigator

On-Call: Application Developme...
Change - All
Incident - All
User Administration - Users
Self-Service - Business Applicati...
IBM Cloud Pak for Watson AIOps
Configuration - Services

① Insert new records or update existing records using an Excel template file. Click Upload to preview the data to be imported. [More Info](#)

< Import external data into Incident

Do you want to insert or update data ?
☒ Insert
☐ Update

Do you want to create an Excel template to enter data ?
☐

Step 1: Upload the template file

Excel template file [Choose file](#) No file chosen

[Upload](#)

5. Click on **Preview Imported Data**

servicenow Service Management

Filter navigator

On-Call: Application Developme...
Change - All
Incident - All
User Administration - Users
Self-Service - Business Applicati...
IBM Cloud Pak for Watson AIOps
Configuration - Services

< Import external data into Incident

Do you want to insert or update data ?
☒ Insert
☐ Update

Do you want to create an Excel template to enter data ?
☐

Step 1: Upload the template file

Excel template file [Choose file](#) import_incidents.xlsx

[Upload](#)

Upload Progress

▶ Importing Excel Workbook Succeeded 100%

Import complete - Succeeded in 5 Seconds

[Preview Imported Data](#)

6. Click on **Complete Import** (if there are errors or warnings just ignore them and import anyway)

[Back](#) [Complete Import](#)

3. Import the Incidents from ./doc/servicenow/import_incidents.xlsx

1. Select **Incidents - All** from the right-hand menu
2. Proceed as for the Changes but for Incidents

4. Now you can finish configuring your Service Now Dev instance with CP4WAIOPS by [going back](#) and continue where you left off at **Testing the ServiceNow Integration**.

8. Some Polishing

8.1 Add LDAP Logins to CP4WAIOPS

- Go to **AI Manager** Dashboard
- Click on the top left "Hamburger" menu
- Select **User Management**
- Select **User Groups** Tab
- Click **New User Group**
- Enter demo (or whatever you like)
- Click Next
- Select **LDAP Groups**
- Search for **demo**
- Select **cn=demo,ou=Groups,dc=ibm,dc=com**
- Click Next
- Select Roles (I use Administrator for the demo environment)
- Click Next
- Click Create

8.2 Monitor Kafka Topics

At any moment you can run **./tools/22_monitor_kafka.sh** this allows you to:

- List all Kafka Topics
- Monitor Derived Stories
- Monitor any specific Topic

8.3 Monitor ElasticSearch Indexes

At any moment you can run `./tools/28_access_elastic.sh` in a separate terminal window.

This allows you to access ElasticSearch and gives you:

- ES User
- ES Password

```
*****
***** AI OPS DEBUG - Enable ElasticSearch remote access *****
*****
Initializing.....
Getting credentials
*****
Already on project "cp4waiops" on server "https://c100-e.eu-de.containers.cloud.ibm.com:30783".
OK

Checking credentials
OK - Elasticsearch Username
OK - Elasticsearch Password

ElasticSearch Access
*****
URL : https://localhost:9200
User : cp4waiops-cartridge
Password : s29tRm1TWa

You can use any ElasticSearch Browser. I usually use https://elasticvue.com/

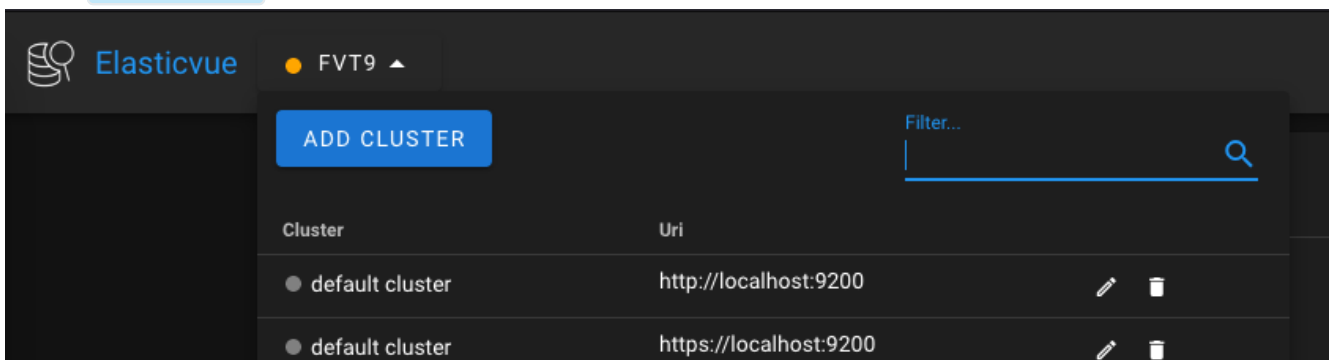
*****
Starting Port Forwarding
*****
Forwarding from 127.0.0.1:9200 -> 9200
Forwarding from [::1]:9200 -> 9200
```

8.3.1 Monitor ElasticSearch Indexes from Firefox

I use the [Elasticvue](#) Firefox plugin.

Follow these steps to connects from Elasticvue:

- Select **Add Cluster**



- Put in the credentials and make sure you put **https** and not **http** in the URL

Add elasticsearch instance

Cluster name
default cluster

Username (optional)
cp4waiops-cartridge

Password (optional)
.....

Uri
https://localhost:9200

Your cluster uses ssl. Make sure that your browser trusts the certificate that you are using, otherwise you will not be able to connect. [Help](#)

TEST CONNECTIONCONNECTCANCEL

- Click **Test Connection** - you will get an error
- Click on the **https://localhost:9200** URL

Add elasticsearch instance


Cluster name
default cluster

Username (optional)
cp4waiops-cartridge

Password (optional)
.....


Uri
https://localhost:9200

Your cluster uses ssl. Make sure that your browser trusts the certificate that you are using, otherwise you will not be able to connect. [Help](#)

 Could not connect. Please make sure that:
1. Your cluster is reachable via <https://localhost:9200>
2. You added the correct settings to your **elasticsearch.yml** and restarted your cluster
Either your cluster is not reachable or you did not configure CORS correctly.

TEST CONNECTIONCONNECTCANCEL

- This will open a new Tab, select **Accept Risk and Continue**



Warning: Potential Security Risk Ahead

Nightly detected a potential security threat and did not continue to localhost. If you visit this site, attackers could try to steal information like your passwords, emails, or credit card details.

What can you do about it?

The issue is most likely with the website, and there is nothing you can do to resolve it.

If you are on a corporate network or using anti-virus software, you can reach out to the support teams for assistance. You can also notify the website's administrator about the problem.

[Learn more...](#)

[Go Back \(Recommended\)](#) [Advanced...](#)

Someone could be trying to impersonate the site and you should not continue.

Websites prove their identity via certificates. Nightly does not trust localhost:9200 because its certificate issuer is unknown, the certificate is self-signed, or the server is not sending the correct intermediate certificates.

Error code: **SEC_ERROR_UNKNOWN_ISSUER**

[View Certificate](#)

[Go Back \(Recommended\)](#) [Accept the Risk and Continue](#)

- Cancel the login screen and go back to the previous tab
- Click **Connect**
- You should now be connected to your AI Manager Elasticsearch instance

Elasticvue default cluster

HOME NODES INDICES SEARCH REST SNAPSHOTS UTILITIES

Indices

NEW INDEX

Show hidden indices

Name	Health	Status	UUID	Aliases	Shards	Lucene docs	Storage		
.1000-1000-20210505-logtrain	yellow	open	MBSnq-510H8MU2ic311zw	[]	1 / 1	315052	109 MB		
.1000-1000-20210506-logtrain	yellow	open	D21np0od556zchK9mVyg	[]	1 / 1	386526	132 MB		
.1000-1000-20211109-logtrain	yellow	open	8HwCOT0RQ7ruacs31f0Aa	[]	8 / 2	24140	8.12 MB		
.1000-1000-20211110-logtrain	yellow	open	fZLYcuc28MECVftuzmBg	[]	8 / 2	82633	33.3 MB		
.1000-1000-changerisk_model_latest	yellow	open	P6hIIVsrdny13Kxfv1A	[]	1 / 1	1	3.84 KB		
.1000-1000-incident_model_latest	yellow	open	-5V2PmavSeSHjyRtVJ7Q	[]	1 / 1	1	3.84 KB		
.1000-1000-lad_registration	yellow	open	RuMa2QWCS_OuBZxQm70Wg	[]	5 / 1	1	277 KB		
.1000-1000-log_model_latest	yellow	open	KSc4j3oLSv28ER0ZDK9WQ	[]	1 / 1	1	3.84 KB		
.1000-1000-usb_model_update	yellow	open	YgKxMwTQ7uQInclqnvW	[]	5 / 1	1	3.95 KB		
.1000-1000-reference_embedding	yellow	open	MjJpatEETrqnaX9f57t0dA	[]	5 / 1	32	127 KB		
.1000-1000-reference_sob	yellow	open	1p0M0eB-RM0hdhQzVt10	[]	5 / 1	32	127 KB		
.1000-1000-v1_model_latest	yellow	open	fAobnrlvSS0KLYaCBV12A	[]	1 / 1	1	4.24 KB		
.1000-1000-v1-anomalies	yellow	open	qpy0d5y0B0dAw9UX3vA	[]	5 / 1	269	1.25 MB		
.1000-1000-v1-applications	yellow	open	rf_L4ELg0LlPMcLm_sH0Q	[]	5 / 1	3	8.78 KB		
.1000-1000-v1-embedding_pca_fe	yellow	open	8JhQX1arPMQczgMsc10NA	[]	5 / 1	3	17.6 KB		
.1000-1000-v1-embedding_pca_model	yellow	open	8Lvl8AL7Q-217r0-Bae0Bw	[]	5 / 1	3	78.5 KB		
.1000-1000-v1-pca_anomaly_group_id	yellow	open	1P2B7R0G0iFqHv1y13ag	[]	5 / 1	4	68.5 KB		
.1000-1000-v1-pca_fe	yellow	open	nAsK3uZ07uB0VY_e4Yfg	[]	5 / 1	3	15.7 KB		
.1000-1000-v1-pca_model	yellow	open	Mkd1lvcCTK7P1R748yV50	[]	5 / 1	3	28 KB		

Rows per page: 100 1-41 of 41

9. Demo the Solution


9.1 Simulate incident

Make sure you are logged-in to the Kubernetes Cluster first

In the terminal type

```
./tools/01_demo/incident_robotshop.sh
```

This will delete all existing Alerts and inject pre-canned event and logs to create a story.

 Give it a minute or two for all events and anomalies to arrive in Slack.

10. TROUBLESHOOTING

10.1 Check with script

! There is a new script that can help you automate some common problems in your CP4WAIOPS installation.

Just run:

```
./tools/10_debug_install.sh
```

and select **Option 1**

10.2 Pods in Crashloop

If the evtmanager-topology-merge and/or evtmanager-ibm-hdm-analytics-dev-inferenceservice are crashlooping, apply the following patches. I have only seen this happen on ROKS.

```
export WAIOPS_NAMESPACE=cp4waiops

oc patch deployment evtmanager-topology-merge -n $WAIOPS_NAMESPACE --patch-file
./yaml/waiops/pazch/topology-merge-patch.yaml

oc patch deployment evtmanager-ibm-hdm-analytics-dev-inferenceservice -n
$WAIOPS_NAMESPACE --patch-file ./yaml/waiops/patch/evtmanager-inferenceservice-
patch.yaml
```


10.3 Pods with Pull Error

If the ir-analytics or cassandra job pods are having pull errors, apply the following patches.

```
export WAIOPS_NAMESPACE=cp4waiops

kubectl patch -n $WAIOPS_NAMESPACE serviceaccount aiops-topology-service-account -p
'{"imagePullSecrets": [{"name": "ibm-entitlement-key"}]}'
kubectl patch -n $WAIOPS_NAMESPACE serviceaccount aiops-ir-analytics-spark-worker -p
'{"imagePullSecrets": [{"name": "ibm-entitlement-key"}]}'
kubectl patch -n $WAIOPS_NAMESPACE serviceaccount aiops-ir-analytics-spark-pipeline-
composer -p '{"imagePullSecrets": [{"name": "ibm-entitlement-key"}]}'
kubectl patch -n $WAIOPS_NAMESPACE serviceaccount aiops-ir-analytics-spark-master -p
'{"imagePullSecrets": [{"name": "ibm-entitlement-key"}]}'
kubectl patch -n $WAIOPS_NAMESPACE serviceaccount aiops-ir-analytics-probablecause -p
'{"imagePullSecrets": [{"name": "ibm-entitlement-key"}]}'
kubectl patch -n $WAIOPS_NAMESPACE serviceaccount aiops-ir-analytics-classifier -p
'{"imagePullSecrets": [{"name": "ibm-entitlement-key"}]}'
kubectl patch -n $WAIOPS_NAMESPACE serviceaccount aiops-ir-lifecycle-eventprocessor-ep -
p '{"imagePullSecrets": [{"name": "ibm-entitlement-key"}]}'
oc delete pod $(oc get po -n $WAIOPS_NAMESPACE | grep ImagePull | awk '{print$1}') -n
$WAIOPS_NAMESPACE
```

10.4 Camel-K Handlers Error

If the scm-handler or snow-handler pods are not coming up, apply the following patches.

```
export WAIOPS_NAMESPACE=cp4waiops

oc patch vaultaccess/ibm-vault-access -p '{"spec":{"token_period":"760h"}}' --type=merge
-n $WAIOPS_NAMESPACE
oc delete pod $(oc get po -n $WAIOPS_NAMESPACE | grep 0/ | grep -v "Completed" | awk
'{print$1}') -n $WAIOPS_NAMESPACE
```

10.5 Slack integration not working

See [here](#)

10.6 Check if data is flowing

10.6.1 Check Log injection

To check if logs are being injected through the demo script:

1. Launch

```
./tools/22_monitor_kafka.sh
```

2. Select option 4

You should see data coming in.

10.6.2 Check Events injection

To check if events are being injected through the demo script:

1. Launch

```
./tools/22_monitor_kafka.sh
```

2. Select option 3

You should see data coming in.

10.6.3 Check Stories being generated

To check if stories are being generated:

1. Launch

```
./tools/22_monitor_kafka.sh
```

2. Select option 2

You should see data being generated.

10.7 Docker Pull secret

! ⚠ Make a copy of the secret before modifying

! ⚠ On ROKS (any version) and before 4.7 you have to restart the worker nodes after the modification

We learnt this the hard way...

```
oc get secret -n openshift-config pull-secret -oyaml > pull-secret_backup.yaml
```

or more elegant

```
oc get Secret -n openshift-config pull-secret -ojson | jq 'del(.metadata.annotations, .metadata.creationTimestamp, .metadata.generation, .metadata.managedFields, .metadata.resourceVersion, .metadata.selfLink, .metadata.uid, .status)' > pull-secret_backup.json
```

In order to avoid errors with Docker Registry pull rate limits, you should add your Docker credentials to the Cluster.

This can occur especially with Rook/Ceph installation.

- Go to Secrets in Namespace **openshift-config**
- Open the **pull-secret** Secret
- Select **Actions / Edit Secret**
- Scroll down and click **Add Credentials**
- Enter your Docker credentials

 Remove Credentials

Registry Server Address *

docker.io

Username *

niklaushirt

Password *

.....



Email

 Add Credentials

Save

Cancel

- Click Save

If you already have Pods in ImagePullBackoff state then just delete them. They will recreate and should pull the image correctly.

11. Uninstall

! The scripts are coming from here <https://github.com/IBM/cp4waiops-samples.git>

If you run into problems check back if there have been some updates.

I have tested those on 3.1.1 as well and it seemed to work (was able to do a complete reinstall afterwards).

Just run:

```
./tools/99_uninstall/3.2/uninstall-cp4waiops.props
```

12. EventManager Installation

12.1. Configure EventManager

! You only have to do this if you have installed EventManager/NOI (As described in Chapter 3.2). For basic demoing with AI Manager this is not needed.

12.1.1 Create Kubernetes Observer for the Demo Applications

This is basically the same as for AI Manager as we need two separate instances of the Topology Manager.

- In the **Event Manager** "Hamburger" Menu select **Administration / Topology Management**
- Under **Observer jobs** click **Configure**
- Click **Add new job**
- Under **Kubernetes**, click on **Configure**
- Choose **local** for **Connection Type**
- Set **Unique ID** to **robot-shop**
- Set **data_center** to **robot-shop**
- Under **Additional Parameters**
- Set **Terminated pods** to **true**
- Set **Correlate** to **true**
- Set Namespace to **robot-shop**
- Under **Job Schedule**
- Set **Time Interval** to 5 Minutes
- Click **Save**

12.1.2 Create REST Observer to Load Topologies

- In the **Event Manager** "Hamburger" Menu select **Administration / Topology Management**
- Under **Observer jobs** click **Configure**
- Click **Add new job**
- Under **REST**, click on **Configure**
- Choose **bulk_replace** for **Job Type**
- Set **Unique ID** to **listenJob** (important!)
- Set **Provider** to **listenJob**
- Click **Save**

12.1.3 Create Merge Rules for Kubernetes Observer

Launch the following:

```
./61_load_robotshop_topology_noi.sh
```

This will create:

- Merge Rules
- Merge Topologies for RobotShop.

! Please manually re-run the Kubernetes Observer to make sure that the merge has been done.

12.1.4 EventManager Webhooks

Create Webhooks in EventManager for Event injection and incident simulation for the Demo.

The demo scripts (in the **demo** folder) give you the possibility to simulate an outage without relying on the integrations with other systems.

At this time it simulates:

- Git push event
- Log Events (Humio)
- Security Events (Falco)
- Instana Events
- Metric Manager Events (Predictive)
- Turbonomic Events
- CP4MCM Synthetic Selenium Test Events

12.1.4.1 Generic Demo Webhook

You have to define the following Webhook in EventManager (NOI):

- **Administration / Integration with other Systems**
- **Incoming / New Integration**
- **Webhook**
- Name it **Demo Generic**
- Jot down the WebHook URL and copy it to the **NETCOOL_WEBHOOK_GENERIC** in the **./tools/01_demo/incident_robotshop-noi.sh** file
- Click on **Optional event attributes**
- Scroll down and click on the + sign for **URL**
- Click **Confirm Selections**

Use this json:

```
{
  "timestamp": "1619706828000",
  "severity": "Critical",
  "summary": "Test Event",
  "nodename": "productpage-v1",
  "alertgroup": "robotshop",
  "url": "https://pirsoscom.github.io/grafana-robotshop.html"
}
```

Fill out the following fields and save:

- Severity: **severity**
- Summary: **summary**
- Resource name: **nodename**
- Event type: **alertgroup**
- Url: **url**
- Description: **"URL"**

Optionnally you can also add **Expiry Time** from **Optional event attributes** and set it to a convenient number of seconds (just make sure that you have time to run the demo before they expire).

12.1.5 Create custom Filter and View in EventManager/ (optional)

12.1.5.1 Filter

Duplicate the **Default** filter and set to global.

- Name: AIOPS
- Logic: **Any** (!)
- Filter:
 - AlertGroup = 'CEACorrelationKeyParent'
 - AlertGroup = 'robot-shop'

12.1.5.2 View

Duplicate the **Example_IBM_CloudAnalytics** View and set to global.

- Name: AIOPS

Configure to your likings.

12.1.6 Create Templates for Topology Grouping (optional)

This gives you probable cause and is not strictly needed if you don't show EventManager!

- In the EventManager "Hamburger" Menu select **Operate / Topology Viewer**
- Then, in the top right corner, click on the icon with the three squares (just right of the cog)
- Select **Create a new Template**
- Select **Dynamic Template**

Create a template for RobotShop:

- Search for **web-deployment** (deployment)
- Create Topology 3 Levels
- Name the template (robotshop)
- Select **Namespace** in **Group type**
- Enter **robotshop_** for **Name prefix**
- Select **Application**
- Add tag **namespace:robot-shop**
- Save

12.1.7 Create grouping Policy

- NetCool Web Gui --> **Insights / Scope Based Grouping**
- Click **Create Policy**
- **Action** select field **Alert Group**
- Toggle **Enabled** to **On**
- Save

12.1.8 Create EventManager/NOI Menu item - Open URL

in the Netcool WebGUI

- Go to **Administration** / **Tool Configuration**
- Click on **LaunchRunbook**
- Copy it (the middle button with the two sheets)
- Name it **Launch URL**
- Replace the Script Command with the following code

```
var urlId = '{$selected_rows.URL}';

if (urlId == '') {
    alert('This event is not linked to an URL');
} else {
    var wnd = window.open(urlId, '_blank');
}
```

- Save

Then

- Go to **Administration** / **Menu Configuration**
- Select **alerts**
- Click on **Modify**
- Move Launch URL to the right column
- Save

12.2 Configure Runbooks

12.2.1 Create Bastion Server

A simple Pod with the needed tools (oc, kubectl) being used as a bastion host for Runbook Automation should already have been created by the install script.

12.2.2 Create the EventManager/NOI Integration

12.2.2.1 In EventManager/NOI

- Go to **Administration** / **Integration with other Systems** / **Automation Type** / **Script**
- Copy the SSH KEY

12.2.2.2 Adapt SSL Certificate in Bastion Host Deployment.

- Select the **bastion-host** Deployment in Namespace **default**
- Adapt Environment Variable SSH_KEY with the key you have copied above.

12.2.3 Create Automation

12.2.3.1 Connect to Cluster

Automation / **Runbooks** / **Automations** / **New Automation**

```
oc login --token=$token --server=$ocp_url --insecure-skip-tls-verify
```

Use these default values

```
target: bastion-host-service.default.svc
user:   root
$token  : Token from your login (from ./tools/20_get_logins.sh)
$ocp_url : URL from your login (from ./tools/20_get_logins.sh, something like
https://c102-e.eu-de.containers.cloud.ibm.com:32236)
```

12.2.3.2 RobotShop Mitigate MySql

[Automation](#) / [Runbooks](#) / [Automations](#) / [New Automation](#)

```
oc scale deployment --replicas=1 -n robot-shop ratings
oc delete pod -n robot-shop $(oc get po -n robot-shop|grep ratings |awk '{print$1}') --
force --grace-period=0
```

Use these default values

```
target: bastion-host-service.default.svc
user:   root
```

12.2.4 Create Runbooks

- [Library](#) / [New Runbook](#)
- Name it [Mitigate RobotShop Problem](#)
- [Add Automated Step](#)
- Add [Connect to Cluster](#)
- Select [Use default value](#) for all parameters
- Then [RobotShop Mitigate Ratings](#)
- Select [Use default value](#) for all parameters
- Click [Publish](#)

12.2.5 Add Runbook Triggers

- [Triggers](#) / [New Trigger](#)
- Name and Description: [Mitigate RobotShop Problem](#)
- Conditions
 - Name: RobotShop
 - Attribute: Node
 - Operator: Equals
 - Value: mysql-instana or mysql-predictive
- Click [Run Test](#)
- You should get an Event [\[Instana\] Robotshop available replicas is less than desired replicas - Check conditions and error events - ratings](#)
- Select [Mitigate RobotShop Problem](#)
- Click [Select This Runbook](#)
- Toggle [Execution](#) / [Automatic](#) to [off](#)
- Click [Save](#)

13. Installing Turbonomic

 Either position the option in the configuratoion file or follow the steps described in this chapter.

13.1 Installing Turbonomic

You can install Turbonomic into the same cluster as CP4WAIOPS.

! You need a license in order to use Turbonomic.

1. Launch

```
ansible-playbook ./ansible/20_install-turbonomic.yaml
```

2. Wait for the pods to come up

3. Open Turbonomic

4. Enter the license

5. Add the default target (local Kubernetes cluster is already instrumented with **kubeturbo**)

It can take several hours for the Supply Chain to populate, so be patient.

13.2 Installing kubeturbo

In order to get other Kubernetes clusters to show up in Turbonomic, you have to install **kubeturbo** (your main cluster is already registered).

1. Adapt `./ansible/templates/kubeturbo/my_kubeturbo_instance_cr.yaml` with the Turbonomic URL and the login

2. Launch

```
ansible-playbook ./ansible/20_1_aiops-addons-kubeturbo.yaml
```

13.3 Turbo to WAIOPS Gateway

! This is not an officially supported tool by any means and is still under heavy development!

In order to push Turbonomic Actions into EventManager you can use my tool.

This tool needs existing **Business Applications**, you can either integrate with Instana (or other APMs) or create one under Settings/Topology.

1. Adapt the `./ansible/templates/turbo-gateway/create-turbo-gateway.yaml` file

Variable	Default Value	Description
POLLING_INTERVAL	'300'	Poll every X seconds
EVTMGR_SUMMARY_PREFIX	'[Turbonomic] '	Prefix in the event summary
EVTMGR_WEBHOOK_URL	netcool-evtmanager.apps.clustername.domain	EventManager hostname
EVTMGR_WEBHOOK_PATH	/norml/xxxx	Webhook URL from EventManager (does not include the hostname, only <code>/norml/xxxx</code>)
TURBO_API_URL	api-turbonomic.apps.clustername.domain	Turbonomic API URL
TURBO_BA_NAME	'RobotShop:robot-shop'	Turbonomic application name in the format APPNAME:ALERTGROUP. This links an EventManager alertgroup with an application
ACTION_STATES	'SUCCEEDED,FAILED,READY,IN_PROGRESS'	The list of ACTION_STATES to filter on
ACTION_TYPES	'MOVE,RESIZE_FOR_PERFORMANCE,RESIZE_FOR EFFICIENCY,RESIZE'	The list of ACTION_TYPES to filter on
DEBUG_ENABLED	'false'	Enable additional log output
ENTITY_TYPES	'VirtualMachine,Application,PhysicalMachine,ContainerSpec,WorkloadController,Container'	The list of ENTITY_TYPES to filter on
ACTION_START_TIME	'-30m'	Period of time in which actions are retrieved. E.g. -5m, -30m, -1h, -1d, -3d, -7d

2. Create Turbonomic Credentials Secret

You can either:

- create the secret from the command line (which will throw a warning for the already existing Secret when installing)

```
oc -n default create secret generic turbo-creds --from-literal=TURBO_USER=<youruser> --from-literal=TURBO_PWD=<yourpw>
```

- replace the secret in the yaml file with

```
oc -n default create secret generic turbo-creds --from-  
literal=TURBO_USER=apiuser --from-literal=TURBO_PWD=turboadmin -o yaml --dry-  
run=client
```

3. Create Generic Webhook in EventManager/NOI with:

```
{  
  "timestamp": "1619706828000",  
  "severity": "Critical",  
  "summary": "Test Event",  
  "nodename": "productpage-v1",  
  "alertgroup": "robotshop",  
  "url": "https://myturbo/something.html"  
}
```

4. Launch

```
ansible-playbook ./ansible/20_3_aiops-addons-turbonomic-gateway.yaml
```


13.4 Generate Metrics

! This is not an officially supported tool by any means and is still under heavy development!

If you have manually created a **Business Applications** you won't get any ResponseTime and Transactions metrics.

With this tool you can add randomized ResponseTime and Transactions metrics to the **Business Application** through the **Data Integration Framework (DIF)**.

Note: The metrics pod can also serve metrics for other **Entity** types (businessApplication, businessTransaction, service, databaseServer, application)

Note: There is also a Route being created by the installer, so that you can test the URLs.

1. Launch

```
ansible-playbook ./ansible/20_2_aiops-addons-turbonomic-metrics.yaml
```

2. Wait for the Pod to become available

3. Add the DIF Target

2. Go to **Settings/Target Configurations**
3. Click **New Target**
4. Select **Custom/DataIngestionFramework**
5. Put in the URL for the metrics (see below) and a name
6. Click **Add**
7. Make sure that Target is green and reads **Validated**

It takes some time for the metrics to start showing up. Polling is every 10 minutes

13.4.1 Test URL

You can use the following URL to test if everything is working:

```
http://turbo-dif-service.default:3000/helloworld
```

This will create a standalone **Business Application** called **Hello World** without any other **Entities** attached to it.

But with metrics being ingested.

13.4.2 Construct the URL

The URL has the format of:

```
http://turbo-dif-service.default:3000/<TYPE>/<NAME>/<UUID>
```

where:

- TYPE: Type of the **Entity**
(businessApplication/businessTransaction/service/databaseServer/application)

- NAME: The name of the **Entity**
- UUID: The UUID that you can find under **Entity Information / Show All / Vendor ID**

So an example might be:

`http://turbo-dif-service.default:3000/service/Service-robot-shop%2Fcatalogue/b2d6fd52-c895-469e-bb98-2a791faefce7`

`http://turbo-dif-service.default:3000/businessApplication/RobotShop/285333133684640`

`http://turbo-dif-service.default:3000/businessTransaction/RobotShopCatalog/285333292080720`

`http://turbo-dif-`

`service.default:3000/businessTransaction/RobotShopPayment/285333133684640`

`http://turbo-dif-service.default:3000/businessTransaction/RobotShopUser/285333133684640`

14. Installing OCP ELK

 Either position the option in the configuratoion file or follow the steps described in this chapter.

You can easily install ELK into the same cluster as CP4WAIOPS.

1. Launch

```
ansible-playbook ./ansible/22_install-elk-ocp.yaml
```

2. Wait for the pods to come up
3. Open Kibana

15. HUMIO

 Either position the option in the configuratoion file or follow the steps described in this chapter.

 This demo supports pre-canned events and logs, so you don't need to install and configure Humio unless you want to do a live integration (only partially covered in this document).

15.1 Create Licence Secret

Before starting the installation you have to create the Secret with the licence information.

```
oc create ns humio-logging
oc create secret generic humio-license -n humio-logging --from-
literal=data=eyJhbGciOiJFUzxxxxaWmdRTrr_ksdfaa
```

15.2 Install Humio and Fluentbit

Just launch the following and this should automatically install:

- Kafka
- Zookeeper
- Humio Core
- Humio Repository
- Humio Ingest Token
- Fluentbit

```
ansible-playbook ./ansible/21_install-humio.yaml
```

15.3 Live Humio integration with AIManager

15.3.1 Humio URL

- Get the Humio Base URL from your browser
- Add at the end `/api/v1/repositories/aiops/query`

15.3.2 Accounts Token

Get it from Humio --> **owl** in the top right corner / **Your Account** / **API Token**

15.3.3 Create Humio Log Integration

- In the **AI Manager** "Hamburger" Menu select **Operate** / **Data and tool integrations**
- Under **Humio**, click on **Add Connection**
- Click **Connect**
- Name it **Humio**
- Paste the URL from above (**Humio service URL**)
- Paste the Token from above (**API key**)
- In **Filters (optional)** put the following:

```
"kubernetes.namespace_name" = /robot-shop/  
| "kubernetes.container_name" != load
```

- Click **Next**
- Put in the following mapping:

```
{  
  "codec": "humio",  
  "message_field": "@rawstring",  
  "log_entity_types": "clusterName, kubernetes.container_image_id, kubernetes.host,  
kubernetes.container_name, kubernetes.pod_name",  
  "instance_id_field": "kubernetes.container_name",  
  "rolling_time": 10,  
  "timestamp_field": "@timestamp"  
}
```

- Click **Test Connection**
- Switch **Data Flow** to the **ON** position **!**
- Select the option for your use case:
 - **Live data for continuous AI training and anomaly detection** if you want to enable log anomaly detection
 - **Live data for initial AI training** if you want to start ingesting live data for later training
 - **Historical data for initial AI training** if you want to ingest historical data to start training rapidly
- Click **Done**

15.3.4 Create Humio Events Integration

Events integration is done via EventManager/NOI.

For the time being this only takes the first alert being pushed over (no way to handle arrays).

The native Humio integration seems to have a bug that gives "mergeAdvanced is not a function".

15.3.4.1 On the EventManager/NOI side

Create a Webhook integration:

Field	Value
Severity	"Critical"
Summary	alert.name
Resource name	events[0].kubernetes.container_name
Event type	events[0].kubernetes.namespace_name

With this sample payload:

```
{
  "repository": "aiops",
  "timestamp": "2021-11-19T15:50:04.958Z",
  "alert": {
    "name": "test1",
    "description": "",
    "query": {
      "queryString": "\"kubernetes.container_name\" = ratings\n| @rawstring = /error/i",
      "end": "now",
      "start": "2s"
    },
    "notifierID": "Rq4a9KUbomSIBvEcdC7kzzmdBtPI3yPb",
    "id": "rCA2w5zaIE6Xr3RKlFfhAxqqbGqGxGLC"
  },
  "warnings": "",
  "events": [
    {
      "kubernetes.annotations.openshift_io/scc": "anyuid",
      "kubernetes.annotations.k8s_v1_cni_cncf_io/network-status": "[{\n  \"name\": \"k8s-pod-network\", \n  \"ips\": [\n    \"172.30.30.153\" \n  ], \n  \"default\": true, \n  \"dns\": { }\n}]",
      "kubernetes.annotations.cni_projectcalico_org/podIPs": "172.30.30.153/32",
      "@timestamp.nanos": "0",
      "kubernetes.annotations.k8s_v1_cni_cncf_io/networks-status": "[{\n  \"name\": \"k8s-pod-network\", \n  \"ips\": [\n    \"172.30.30.153\" \n  ], \n  \"default\": true, \n  \"dns\": { }\n}]",
      "kubernetes.pod_name": "ratings-5d9dff56bd-864kq",
      "kubernetes.labels.service": "ratings",
      "kubernetes.annotations.cni_projectcalico_org/podIP": "172.30.30.153/32",
      "kubernetes.host": "10.112.243.226",
      "kubernetes.container_name": "ratings",
      "kubernetes.labels.pod-template-hash": "5d9dff56bd",
    }
  ]
}
```

```

    "kubernetes.docker_id":
"87a98617a14684c02d9d52a6245af377f8b1a246d196f232cad494a7a2d125b7",
    "@ingesttimestamp": "1637337004272",
    "kubernetes.container_hash": "docker.io/robotshop/rs-
ratings@sha256:4899c686c249464783663342620425dc8c75a5d59ca55c247cf6aec62a5ffff1a",
    "kubernetes.container_image": "docker.io/robotshop/rs-ratings:latest",
    "#repo": "aiops",
    "@timestamp": 1637337003872,
    "kubernetes.namespace_name": "robot-shop",
    "@timezone": "Z",
    "@rawstring": "2021-11-19T09:50:03.872288692-06:00 stdout F [2021-11-19 15:50:03]
php.INFO: User Deprecated: Since symfony/http-kernel 5.3:
\\\"Symfony\\Component\\HttpKernel\\Event\\KernelEvent::isMasterRequest()\\\" is deprecated,
use \\\"isMainRequest()\\\" instead. {\\\"exception\\\":\\\"[object] (ErrorException(code: 0):
User Deprecated: Since symfony/http-kernel 5.3:
\\\"\\\"Symfony\\Component\\HttpKernel\\Event\\KernelEvent::isMasterRequest()\\\"\\\" is
deprecated, use \\\"\\\"isMainRequest()\\\"\\\" instead. at /var/www/html/vendor/symfony/http-
kernel/Event/KernelEvent.php:88)\\\"} [ ]",
    "@id": "tiMU0F8kdNf6x0qMduS9T31q_269_400_1637337003",
    "kubernetes.pod_id": "09d64ec8-c09f-4650-871f-adde27ca863e",
    "#type": "unparsed",
    "kubernetes.annotations.cni_projectcalico_org/containerID":
"337bf300371c84500756a6e94e58b2d8ee54a1b9d1bc7e38eb410f1c1bbd6991"
  }
],
"numberOfEvents": 1
}

```

15.3.4.2 On Humio:

- Create Action:
 - Use the Webhook from EventManager/NOI
 - Select **Skip Certificate Validation**
 - Click **Test Action** and check that you get it in EventManager/NOI Events
- Create Alert:
 - With Query (for example):

```

"kubernetes.container_name" = ratings
| @rawstring = /error/i

```

- Time Window 2 seconds
- 1 second throttle window
- Add action from above

16. ServiceMesh

 Either position the option in the configuratoion file or follow the steps described in this chapter.

You can easily install ServiceMesh/Istio into the same cluster as CP4WAIOPS.

This will instrument the RobotShop Application at the same time.

1. Launch

```
ansible-playbook ./ansible/29_install-servicemesh.yaml
```

2. Wait for the pods to come up

3. You can get the different URLs (RobotShop, Kibana, Grafana, Jaeger) by launching:

```
./tools/20_get_logins.sh > my_credentials.txt
```

17. AWX

 Either position the option in the configuratoion file or follow the steps described in this chapter.

You can easily install AWX (OpenSource Ansible Tower) into the same cluster as CP4WAIOPS.

1. Launch

```
ansible-playbook ./ansible/23_install-awx.yaml
```

2. Wait for the pods to come up

3. You can get the URLs and access credentials by launching:

```
./tools/20_get_logins.sh > my_credentials.txt
```

17.1. Configure AWX

There is some demo content available to RobotShop.

1. Log in to AWX

2. Add a new Project

1. Name it **DemoCP4WAIOPS**
2. Source Control Credential Type to **Git**
3. Set source control URL to **https://github.com/niklaushirt/ansible-demo**
4. Save

3. Add new Job Template

1. Name it **Mitigate Robotshop Ratings Outage**
2. Select Inventory **Demo Inventory**
3. Select Project **DemoCP4WAIOPS**
4. Select Playbook **cp4waiops/robotshop-restart/start-ratings.yaml**
5. Select **Prompt on launch** for **Variables** **!**
6. Save

17.2. Configure AWX Integration

In EventManager:

1. Select **Administration** / **Integration with other Systems**
2. Select **Automation type** tab
3. For **Ansible Tower** click **Configure**
4. Enter the URL and credentials for your AWX instance (you can use the default **admin** user)
5. Click Save

17.3. Configure Runbook

In EventManager:

1. Select **Automations** / **Runbooks**
2. Select **Library** tab
3. Click **New Runbook**
4. Name it **Mitigate Robotshop Ratings Outage**
5. Click **Add automated Step**
6. Select the **Mitigate Robotshop Ratings Outage** Job
7. Click **Select this automation**
8. Select **New Runbook Parameter**
9. Name it **ClusterCredentials**
10. Input the login credentials in JSON Format (get the URL and token from the 20_get_logins.sh script)

```
{
  "my_k8s_apiurl": "https://c117-e.xyz.containers.cloud.ibm.com:12345",
  "my_k8s_apikey": "PASTE YOUR API KEY"
}
```

11. Click Save
12. Click Publish

Now you can test the Runbook by clicking on **Run**.

